

Ambiguity Resolution Brief

Student: ALIA ISLAM

Date: June 11, 2026

1. Founder Request

The founder wants recruiters to be able to manually override AI-generated candidate scores and shortlist candidates themselves before the pilot launches. The request was delivered as a single urgent sentence with no definition of scope, permissions, audit requirements, or fairness constraints.

2. Known Facts

Known Fact	Source
HireSignal uses AI to score and rank candidates	Scenario Package
A recruiter pilot is imminent	Founder message
Recruiters want to shortlist candidates the AI scored low	Recruiter message
Engineering is concerned about unlogged, unexplained overrides	Engineering Note
No full admin permissions system exists yet	Business Constraints
No big redesign is expected before the pilot	Business Constraints
A minimal safe v1 is acceptable to stakeholders	Business Constraints

3. Open Ambiguities

Ambiguity	Why It Matters	Risk If Ignored
Who is allowed to override? All recruiters or only senior/designated ones?	Determines permission model needed before any code is written	Any recruiter overrides any candidate silently, trust collapse if discovered
Is a reason required when overriding?	Decides whether a free-text field is needed or just a button	Unexplained overrides look arbitrary to candidates and auditors; fairness liability
Should overrides be logged/audited?	Changes backend requirements significantly	No traceability, cannot investigate bias or abuse after the fact
Should the original AI score remain visible after override?	Affects UI and whether AI trust is preserved or hidden	Hiding the score removes transparency; showing it may confuse recruiters
Is an override reversible?	Changes state machine design of candidate status	Permanent override with no reversal = accidental shortlisting becomes irreversible

4. One Clarification Question

Question asked:

When a recruiter overrides an AI score to shortlist a candidate, should a reason be required or is a silent one-click override acceptable for the pilot?

Why this was the highest-leverage question:

This single question collapses the biggest cluster of risk. The answer determines:

- Whether any audit/logging infrastructure is needed
- Whether the feature is legally and ethically safe to ship
- Whether the UI needs a reason input field or just a button
- Whether the pilot will create fairness exposure

Every other ambiguity (who can override, is it logged, is it reversible) becomes easier to scope once we know if reason capture is required.

5. Bounded Assumptions

These assumptions apply if the founder confirms a minimal v1 approach.

Assumption	Why It Is Reasonable	Risk If Wrong
All pilot recruiters have equal override permissions	Pilot has a small, trusted recruiter group, no need for role hierarchy yet	If some recruiters shouldn't override, we've shipped uncontrolled access
A short reason field (optional for v1) is sufficient for audit safety	Captures intent without blocking action; can be made mandatory post-pilot	Blank reasons still leave audit gaps; legal/compliance risk if override is challenged
AI score remains visible alongside the manual shortlist status	Preserves transparency; low implementation cost	Recruiter confusion if both statuses are shown without clear labeling
Override is reversible by the same recruiter	Prevents accidents; no approval workflow needed for pilot scale	If reversal is uncontrolled, shortlist status becomes unreliable
Override is logged server-side even if not shown in UI	Minimal backend effort; creates safety net for later review	If logging is skipped entirely, there is no recovery path for disputes

6. Minimal Safe V1

V1 Behavior	Why It Is Safe Enough
A "Shortlist" button on each candidate card	Directly answers the recruiter's need; minimal UI surface
Optional free-text reason field (1–2 sentences)	Captures intent; doesn't block speed; upgradeable to required later
Server-side log of: recruiter ID, candidate ID, timestamp, reason (if entered)	Invisible to UI but creates audit trail with zero design cost
AI score remains visible but marked as "AI Score" vs "Manually Shortlisted"	Transparency without confusion; no data is hidden

V1 Behavior	Why It Is Safe Enough
Recruiter can undo their own shortlist	Prevents irreversible mistakes; no approval chain needed at pilot scale

7. What Not To Build Yet

Not In V1	Why Not Now	Risk Of Building It Now
Role-based permissions system	No multi-tier recruiter hierarchy in pilot; overkill	Delays pilot, adds complexity, may be wrong hierarchy anyway
Approval workflow (manager approves override)	Pilot is testing recruiter control, not approval chains	Adds friction that defeats the purpose of giving recruiters control
Audit dashboard / override analytics UI	Logging is sufficient for pilot; a dashboard needs design time	Sinks time into reporting no one will use during a 2-week pilot
Bulk override	Not requested; high fairness risk without individual reasoning	Could allow mass arbitrary shortlisting with no accountability
Candidate notification on override	Premature; candidate-facing features need separate scoping	Wrong message to candidates; creates expectations before product is ready
Full admin override management module (AI suggestion)	Solves a problem we don't have yet	Delays pilot, hides the real ambiguity behind feature volume

8. Minimal Interface Expectation:

Action: Recruiter clicks "Shortlist" on a candidate card

Required fields: Reason text (optional for v1, free text, max 200 chars)

Visible result: Candidate card shows "Manually Shortlisted" badge

Original AI score remains visible, labeled "AI Score"

Error/blocked: If server log fails, action is blocked with message:

"Could not save shortlist. Please try again."

Audit/risk note: Server logs recruiter ID, candidate ID, timestamp,

reason (null if not entered). Not shown in UI for v1.

Undo: Recruiter can click "Remove from Shortlist" on same card.

Undo is also logged.

9. AI / Founder Assumption Challenge

Founder assumption challenged:

The founder assumed "add it quickly" means start building immediately. I challenged this by identifying that the request has five open ambiguities, any one of which, if assumed incorrectly, could create a fairness liability or an audit gap that damages recruiter trust during the pilot. Moving quickly on an unclear requirement is slower than pausing for one targeted clarification.

AI recommendation rejected:

AI suggested building a full override management module with role-based permissions, approval queues, audit dashboard, bulk override, candidate notifications, and analytics.

Why rejected:

This recommendation solves problems that don't exist yet at pilot scale, hides the core ambiguity (whether a reason is required) behind feature volume, risks delaying the pilot, and makes assumptions about the recruiter permission hierarchy that haven't been validated. It is the classic overbuild trap, it sounds thorough but does not answer whether v1 is safe to ship at all.

10. Founder-Ready Decision

Decision: Proceed Conditionally

Condition: Founder confirms whether override reason is required or optional for the pilot.

Founder update:

Respected Founder,

I've reviewed the override request. Here's where I landed. The core need is clear: recruiters want to manually shortlist candidates the AI scored low. That's valid and buildable. Before I scope the work, I have one question:

When a recruiter overrides an AI score, should a reason be required or is a silent one-click shortlist acceptable for the pilot?

This single answer determines whether we have an audit risk on day one.

Regardless of your answer, here's what I'm confident is safe for v1:

- A "Shortlist" button per candidate
- Optional reason field (can be made required later)
- Server-side logging (recruiter, candidate, timestamp, reason)
- AI score stays visible alongside shortlist badge
- Recruiter can undo their own shortlist

What I'm not building yet:

- Role-based permissions system
- Manager approval workflow
- Audit dashboard
- Bulk override
- Candidate notifications

This is minimal, safe, and shippable before the pilot. I can move within 24 hours of your response.

Yours Sincerely,
Alia Islam