

Safe Execution Decision

Student: ALIA ISLAM
Variant ID: D4-A
Date: June 11, 2026

Minimal Safe V1

The smallest safe version of this work is:

A per-candidate "Shortlist" button with:

- An optional free-text reason field (max 200 chars)
- A server-side log entry on every shortlist and undo action
- The original AI score remaining visible, labeled distinctly
- A recruiter-controlled undo action on any shortlist they created

This requires no permissions redesign, no new admin roles, no dashboard, and no candidate-facing changes. It can be scoped, built, and reviewed before the pilot without delaying it.

Risks Prevented:

Risk	How This Decision Reduces It
Permission risk	No new permission tiers introduced. All pilot recruiters get equal access. A known, small, trusted group. No unknown actors can override.
Fairness risk	AI score stays visible alongside manual shortlist. Decisions are not hidden. Optional reason field captures intent even when not required.
Audit risk	Server-side log captures recruiter ID, candidate ID, timestamp, and reason on every action including undos. No dashboard needed for the pilot; the data exists if investigated.

Risk	How This Decision Reduces It
Overbuild risk	Full module (permissions, approval queue, bulk override, dashboard, notifications) is explicitly deferred. Pilot scope is locked to one button, one field, one log.

Rejected Scope:

Rejected	Reason
Role-based permissions system	No recruiter hierarchy exists in the pilot; premature architecture
Manager approval workflow	Adds friction that defeats the purpose of recruiter autonomy
Audit dashboard / analytics UI	Logging covers pilot needs; dashboard is post-pilot scope
Bulk override	High fairness risk without individual reasoning; not requested
Candidate notifications	Candidate-facing feature requiring separate scoping and consent design
AI-suggested full override management module	Overbuild; hides the actual ambiguity behind feature volume

Final Judgment:

The founder's request was real but compressed. "Add it quickly" was pressure, not a specification. Proceeding without clarification would have meant building either too little (a button with no audit trail) or too much (a full permissions module that delays the pilot). The safe path was to identify the one question that collapses the most risk whether a reason is required and scope only what can be safely built regardless of the answer.

My decision is Proceed Conditionally: the v1 scope is defined and ready to execute the moment the founder confirms the reason field requirement. Nothing in the v1 creates irreversible decisions. The audit log is silent infrastructure that costs nothing but protects everything. The AI recommendation was rejected not because it was wrong in principle, but because it was wrong for this moment a pilot with a small team, tight timeline, and an unresolved fairness question that volume of features would only obscure.

Good engineering here is knowing what not to build, not just what to build.